

MNIST Digit Classifier – Report (Feedforward Neural Network)

Objective

To train a neural network using **PyTorch** that can accurately classify handwritten digits (0–9) from the MNIST dataset. The model should be evaluated for accuracy and performance using appropriate visualizations.

Model Architecture and Design Choices


- **Model Type:** Feedforward Neural Network (Fully Connected)
- **Input Layer:** 28×28 image flattened to a 784-dimensional vector
- **Hidden Layers:**
 - Layer 1: 128 neurons + ReLU
 - Layer 2: 64 neurons + ReLU
- **Output Layer:** 10 neurons (one for each digit class)
- **Loss Function:** CrossEntropyLoss (suitable for multi-class classification)
- **Optimizer:** Adam with learning rate = 0.001
- **Training Epochs:** 10
- **Batch Size:** 64
- **Data Normalization:** Mean = 0.1307, Std = 0.3081 (MNIST standard)

Key Findings

- The model achieved **97.35% test accuracy** after 10 epochs of training.
- The training loss consistently decreased over epochs, indicating good convergence.
- The confusion matrix showed that the model performed well on all classes, with a few confusions primarily between similar-looking digits like **5 and 3**.
- A simple fully connected model can perform quite well on MNIST, though CNNs may improve performance further.

Plots & Visual Results

◆ Loss vs. Epochs



This plot shows the steady decline in training loss, validating effective learning.

◆ Confusion Matrix

